

715 TO 530

SFT dataset remediation ledger

A single evidence-backed account of the raw reverify corpus, the 72 terminal bad cases, the 129 rewritten verified complements, the retained anchors, and the completed training/evaluation receipts.

715	72	530	129
RAW ROWS AUDITED	HARD REJECTS ACCOUNTED	FINAL TRAINING ROWS	EXECUTABLE-PASSED COMPLEMENTS

Prepared from the frozen 715-row archive, exhaustive audit ledgers, final 530-row package, combined-v5 verification receipts, and Modal run archive.

1. Executive record

The phrase '715 to 530' is shorthand, not a literal filter. The 715-row reverify archive was first concatenated with 160 Gold-v4 rows and trained as an 875-row corpus. The later 530-row corpus was rebuilt from 401 prior curated anchors plus 129 rewritten rows selected from the combined-v5 verification lane.

CLAIM BOUNDARY

The audit establishes defects and missing evidence in the 715-row archive. The final package establishes its own composition and gates. The evaluation numbers are recorded without attributing any metric to a particular row or remediation action.

Area	Measure A	Measure B	Measure C
Raw archive disposition	72 reject	643 review	0 train-ready
Hard-reject mechanisms	40 contract contradictions	25 strict compile failures	7 same-ID Gold conflicts
Final corpus construction	401 inherited anchors	129 verified complements	530 unique task IDs
Raw-to-final identity	176 shared task IDs	176 prompt+answer rewrites	539 raw task IDs absent
Hard-reject outcome	65 excluded	7 trusted ancestors retained	0 raw reject targets retained

What was changed

Control	Raw-concat path	Final-530 path
Mixture	Raw concat: 160 + 715	Final: 401 anchors + 129 complements
Admission	All 715 appended	Complement requires executable_stage=passed and a family absent from the 401 anchors
Lineage	7 conflicting targets coexist by task ID	7 raw targets removed; the prior curated versions remain
Terminal defects	40 contradictions and 25 compile failures included	All 65 task IDs absent from the 530
Review-row promotion	No row-level semantic verification supplied	129 rewritten prompt/target pairs carry verified_reference=true and executable-stage receipts
Duplication and leakage gates	7 duplicate task IDs in the 875 concat	0 duplicate IDs, prompt hashes, answer hashes, or fixed-26 task-ID overlaps

2. The 715-row audit

Gate	Recorded result	Meaning
Rows parsed	715/715	Mechanical integrity
Unique task IDs	715/715	No duplicate ID inside the archive
Exact editable file set	715/715	Assistant returned exactly the supplied files
Supplied API symbols preserved	715/715	Local starter API inventory retained
Within 4,096 tokens	715/715	No token-limit rejection
Strict C++17 compile	690/715	25 terminal failures
Exact fixed-26 row overlap	0	No exact benchmark-row contamination found
Supplied row-level semantic verification	0/715	643 nonterminal rows remained review-only

Why compile was not enough

Compilation proved that 690 targets were syntactically viable under the pinned Apple Clang C++17 warnings-as-errors command. It could not prove behavior, edge cases, contract alignment, or lineage consistency. The 40 contradictory pairs compiled while positively rewarding behavior that disagreed with the prompt's own discriminator.

GOOD EVIDENCE

The audit paired every negative claim with a verified counterexample. Examples include flower-field for transformation-contract alignment, bbq-audio-processing for the bounded queue family, ETL for table transformation, and the same-ID Gold ancestors for the seven lineage conflicts.

3. Final 530 corpus construction

Component	Count/value	Receipt
Original anchors	401	Inherited unchanged from train-401-dsa.zip
Complement rows	129	Selected from combined-v5 reviewed-new rows
Complement family rule	20 families	Family absent from the original 401 set
Complement executable stage	129 passed	All selected rows marked executable_stage=passed
Prompt/target lineage	129/129 changed	Both prompt and target differ from the raw-715 version
Verification marker	129/129 true	verified_reference=true
Mutation score	min 0.7222	median 1.0000; mean 0.9751
Semantic confidence	min 0.90	median 0.98; mean 0.970

Independent review and repair routes

Semantic route	Rows
original-independent-judge	58
targeted-repair-independent-rejudge	42
numeric-domain-rejudge	24
targeted-repair-wave2-independent-xhigh-rejudge	5

V5 EVIDENCE LIMIT

Exploratory fast-close corpus: all rows passed the pinned source compile audit and semantic repair review. New rows are ranked by the latest available executable diagnostics. The prior full aggregate is used where the interrupted hardened replay has no terminal journal; raw mutation outcomes, generated-test compile failures, and generator-uncertain rows are not represented as independently adjudicated mutation passes.

4. Formatting and reviewability profile

Static indicator	Raw 715	Final 530	401 anchors	129 complements
namespace curriculum in target	705/715 (98.6%)	399/530 (75.3%)	399/401	0/129
Target has a line over 180 chars	606/715 (84.8%)	370/530 (69.8%)	370/401	0/129
Duplicate '# Instructions' heading	498/715 (69.7%)	401/530 (75.7%)	401/401	0/129

The 129-row complement contributes none of the three listed static style indicators. The remaining instances in the final corpus come from the inherited 401-row anchor set. These are reviewability indicators, not correctness proofs.

Family distribution of the 129-row complement

Family	Rows	Family	Rows
ascii-art-scaling	16	clock-arithmetic	9
countdown-timers	7	cross-midnight-intervals	5
date-difference	4	duration-formatting	7
elapsed-time-accumulation	4	flood-fill	2
future-date-calculations	4	grid-ownership-mapping	15
leap-year-rule	5	maze-to-graph	2
nth-and-final-occurrences	2	offset-aware-range-overlap	7
overflow-safe-date-math	3	pattern-printing	3
run-length-image-encoding	6	sparse-matrix-encoding	7
table-pivot	17	text-justification	4

5. Training receipt

Field	Recorded value
Run	glm47-aider-complement-530-sft-20260721
Hardware	8 x H100 80GB
Packaged rows	530
Rows consumed	520; final 10-row remainder dropped by the complete batch-20 loader
Optimizer iterations	26
Final loss	0.15497849990254914
Final train throughput	8,110.5166 tokens/s
Peak memory	66,541 MiB
Adapter SHA-256	f1ea45bc327dc6e28d0287aea75c6b691e99d2ec2f7fdb7f07bbbf5ccd6cf36a
Modal accrued cost	\$25.61 for complementary-530 profile, SFT, finalization, and eval app IDs

Fixed-26 evaluation receipts

Metric	Base	Original 401 v1	Raw 160+715	Final 530
Pass@1	0/26 (0.0%)	1/26 (3.8%)	0/26 (0.0%)	0/26 (0.0%)
Pass@2	4/26 (15.4%)	5/26 (19.2%)	3/26 (11.5%)	7/26 (26.9%)
Well-formed	26/26	26/26	26/26	26/26
Total tokens	2,212,419	1,732,287	1,496,776	1,611,304
User asks	164	146	140	158
Seconds/case	339.7	481.4	454.4	518.7
Error outputs	6	9	4	6
Context exhaustions	6	9	4	6

Evaluation configuration for the final 530 receipt: 4 x H100, SGLang TP=4, max-running-requests 16, Aider threads 8, two tries, temperature 0.7, top-p 1.0, and max output 32,768 tokens. Pass@2 is the two-try repair-based benchmark field, not two independent samples.

6. How to read the 72-case ledger

Ledger field	Meaning
BAD EVIDENCE	The concrete compile diagnostic, prompt/target contradiction, or same-ID conflict recorded in the exhaustive audit.
530 TREATMENT	Whether the raw task ID was excluded or replaced by the prior curated ancestor. Exclusion is not mislabeled as a repair.
GOOD REFERENCE	A verified example named by the audit so the negative case teaches the desired positive behavior.
TRAINING RULE	The reusable admission rule extracted from the case.

COMPLETE ACCOUNTING

The following appendix contains all 72 terminal rejects exactly once: 40 contract contradictions, 25 strict compilation failures, and 7 same-ID conflicts. The other 643 rows were review-only, not proven bad; 129 later re-entered as changed prompt/target pairs with verification receipts.

Appendix A1. Prompt/target contradictions (40)

Each row explicitly names a required behavior and a mutation-like discriminator, while the positive target implements the ordinary operation. All 40 raw targets are absent from the final corpus.

01. Raw row 536 - compose-orientation-commands matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: dihedral transform composition state machine. A renamed generic dense-grid rotation is Discriminator: toggle reflection metadata without reflecting the coordinate. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
02. Raw row 537 - rotate-aerial-tiles matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: in-place concentric four-cycle square rotation. A renamed generic dense-grid rotation is Discriminator: break the four-way layer cycle by copying the right edge twice. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
03. Raw row 538 - rotate-block-sparse-layout matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: two-level sparse block permutation and dense micro-rotation. A renamed generic dense-grid rotation is Discriminator: move block coordinates but leave dense blocks unrotated. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
04. Raw row 539 - rotate-directional-glyph-map matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: rectangular byte-grid transform with symbol automaton. A renamed generic dense-grid rotation is Discriminator: leave north arrows unchanged after rotation. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

05. Raw row 540 - rotate-directional-vector-field matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: spatial remap plus signed basis-vector transform. A renamed generic dense-grid rotation is Discriminator: move vector cells without rotating their components. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
06. Raw row 541 - rotate-embedded-route-graph matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: graph identity validation plus coordinate and heading transform. A renamed generic dense-grid rotation is Discriminator: rotate node coordinates without rotating headings. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
07. Raw row 542 - rotate-lazy-grid-view matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: lazy orientation state with inverse query mapping. A renamed generic dense-grid rotation is Discriminator: query the transposed rather than rotated source coordinate. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
08. Raw row 543 - rotate-oriented-tile-atlas matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: object-layout permutation with modular metadata composition. A renamed generic dense-grid rotation is Discriminator: move tiles without composing their orientation. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

09. Raw row 544 - rotate-packed-occupancy-mask matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: direct packed-bit index permutation. A renamed generic dense-grid rotation is Discriminator: copy each occupied bit to its original index. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
10. Raw row 545 - rotate-perimeter-conveyor matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: clockwise perimeter extraction and cyclic shift. A renamed generic dense-grid rotation is Discriminator: leave the perimeter cycle unshifted. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
11. Raw row 546 - rotate-planar-channels matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: multi-plane shape validation and synchronized remap. A renamed generic dense-grid rotation is Discriminator: leave each channel plane unrotated. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
12. Raw row 547 - rotate-polyline-waypoints matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: signed Cartesian point transform with polyline metric. A renamed generic dense-grid rotation is Discriminator: reflect instead of clockwise rotating signed points. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

13. Raw row 548 - rotate-quadtree-leaf-bounds matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: power-of-two alignment validation and Morton-key ordering. A renamed generic dense-grid rotation is Discriminator: retain original quadtree leaf bounds. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
14. Raw row 549 - rotate-rectangular-regions matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: half-open rectangle bound transform. A renamed generic dense-grid rotation is Discriminator: return original rectangle bounds. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
15. Raw row 550 - rotate-rectangular-scan matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: out-of-place rectangular coordinate remap. A renamed generic dense-grid rotation is Discriminator: transpose without reversing the destination column. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
16. Raw row 551 - rotate-row-run-encoding matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: RLE validation, coordinate emission, and canonical recompression. A renamed generic dense-grid rotation is Discriminator: transpose decoded samples without reversing source rows. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

17. Raw row 552 - rotate-sparse-board-coordinates matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: validated sparse coordinate transform with canonical sort. A renamed generic dense-grid rotation is Discriminator: leave sparse coordinates unrotated. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
18. Raw row 553 - rotate-square-subwindow matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: bounded submatrix copy-back with untouched exterior. A renamed generic dense-grid rotation is Discriminator: copy the selected window without rotating it. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
19. Raw row 554 - rotate-stencil-kernel matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: numeric kernel permutation with anchor transform. A renamed generic dense-grid rotation is Discriminator: rotate weights but not the anchor column. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
20. Raw row 555 - rotate-voxel-slice-stack matrix-rotation-90 / grid-matrix-image	BAD EVIDENCE Required mechanism: three-dimensional slice remap with depth reversal. A renamed generic dense-grid rotation is Discriminator: rotate slices without reversing depth order. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

21. Raw row 696 - transpose-anti-diagonal transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: dimension-swapping reverse-row reverse-column coordinate transform. A renamed generic row-to-column loop is Discriminator: perform a main-diagonal transpose instead of anti-diagonal reflection. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
22. Raw row 697 - transpose-banded-matrix transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: band-coordinate reconstruction and transposed compact reindexing. A renamed generic row-to-column loop is Discriminator: negate every copied band value during transposition. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
23. Raw row 698 - transpose-block-sparse transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: two-level block-coordinate and in-block transpose. A renamed generic row-to-column loop is Discriminator: transpose block coordinates but leave each dense block unchanged. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
24. Raw row 699 - transpose-call-center transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: out-of-place dense rectangular coordinate bijection. A renamed generic row-to-column loop is Discriminator: reverse source columns instead of applying the row-to-column bijection. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

25. Raw row 700 - transpose-channel-planes transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: shape validation plus synchronized interleaved-to-planar remap. A renamed generic row-to-column loop is Discriminator: copy blue samples into the green plane. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
26. Raw row 701 - transpose-common-rectangle transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: shortest-row width selection with deterministic tail accounting. A renamed generic row-to-column loop is Discriminator: discard the required ignored-tail count. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
27. Raw row 702 - transpose-coordinate-coalesce transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: validated coordinate swap with ordered-map coalescing. A renamed generic row-to-column loop is Discriminator: coalesce at the original rather than transposed coordinates. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
28. Raw row 703 - transpose-csr-to-csc transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: count-prefix-scatter compressed sparse conversion. A renamed generic row-to-column loop is Discriminator: scatter each nonzero into the following sparse column. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

29. Raw row 704 - transpose-directed-graph transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: edge reversal with duplicate validation and sorted adjacency. A renamed generic row-to-column loop is Discriminator: preserve the original direction of every edge. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
30. Raw row 705 - transpose-lazy-view transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: inverse coordinate mapping over immutable source storage. A renamed generic row-to-column loop is Discriminator: shift every inverse-mapped source column by one. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
31. Raw row 706 - transpose-packed-bitboard transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: direct packed-bit row-column index permutation. A renamed generic row-to-column loop is Discriminator: leave every occupied bit at its original index. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
32. Raw row 707 - transpose-permutation-ledger transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: permutation validation, inverse construction, and visited-cycle decomposition. A renamed generic row-to-column loop is Discriminator: return the original permutation instead of its inverse. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

33. Raw row 708 - transpose-ragged-pad-mask transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: zip-longest ragged transpose with explicit provenance mask. A renamed generic row-to-column loop is Discriminator: erase provenance for every real input cell. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
34. Raw row 709 - transpose-relation-index transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: ordered relation inversion with per-pair duplicate suppression. A renamed generic row-to-column loop is Discriminator: group by the original left endpoint. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
35. Raw row 710 - transpose-run-encoded-image transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: validated run decoding, coordinate transpose, and adjacent-run recompression. A renamed generic row-to-column loop is Discriminator: never merge adjacent equal values in an output row. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
36. Raw row 711 - transpose-square-in-place transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: in-place upper-triangle pair swaps with atomic validation. A renamed generic row-to-column loop is Discriminator: self-swap the upper triangle and leave the matrix unchanged. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

37. Raw row 712 - transpose-stream-batches transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: sequential column batching with explicit continuation offsets. A renamed generic row-to-column loop is Discriminator: advance one column and emit overlapping batches. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
38. Raw row 713 - transpose-symmetric-triangle transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: packed triangular index traversal with mirrored-coordinate validation. A renamed generic row-to-column loop is Discriminator: discard every mirrored off-diagonal value. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
39. Raw row 714 - transpose-tensor-axes transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: validated axis permutation and flat-stride coordinate remap. A renamed generic row-to-column loop is Discriminator: retain flattened source order instead of applying the axis permutation. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.
40. Raw row 715 - transpose-tiled-grid transpose / grid-matrix-image	BAD EVIDENCE Required mechanism: blocked row-column traversal with clipped boundary tiles. A renamed generic row-to-column loop is Discriminator: reverse rows inside every visited tile. The positive target instead implements the normal operation described by the audit.	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The contradictory pair was not treated as a safe direct repair.
PROMPT/TARGET CONTRADICTION	GOOD REFERENCE flower-field: verified grid transformation with explicit contract and current executable Exercism tests	TRAINING RULE Never reward a target that violates an explicit discriminator. Resolve the specification, regenerate the target, and verify the intended behavior before admission.

Appendix A2. Strict compilation failures (25)

These rows fail the pinned strict Apple Clang C++17 warnings-as-errors gate. All 25 task IDs are absent from the final corpus.

41. Raw row 22 - alert-retention-ring nth-and-final-occurrences / temporal-calendar	BAD EVIDENCE misleading indentation; statement is not part of the previous 'for' [-Werror,-Wmisleading-indentation]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE high-scores: verified sequence task with current executable Exercism tests	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
42. Raw row 71 - bbq-audio-processing-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
43. Raw row 72 - bbq-build-worker-pool-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'last_accepted_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
44. Raw row 73 - bbq-customer-support-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
45. Raw row 74 - bbq-database-write-behind-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.

46. Raw row 75 - bbq-document-indexer-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
47. Raw row 76 - bbq-email-delivery-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
48. Raw row 77 - bbq-file-scan-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
49. Raw row 78 - bbq-fraud-review-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
50. Raw row 79 - bbq-image-upload-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
51. Raw row 80 - bbq-log-writer-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.

52. Raw row 81 - bbq-network-message-pump-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
53. Raw row 82 - bbq-notification-delivery-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
54. Raw row 83 - bbq-order-kitchen-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
55. Raw row 84 - bbq-payment-retry-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
56. Raw row 85 - bbq-print-dispatch-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
57. Raw row 86 - bbq-route-calculation-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.

58. Raw row 87 - bbq-sensor-fusion-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
59. Raw row 88 - bbq-telemetry-ingest-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
60. Raw row 89 - bbq-video-transcode-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
61. Raw row 90 - bbq-warehouse-pick-v2 bounded-blocking-queue / queues-buffers-caches-lists	BAD EVIDENCE private field 'submission_ceiling_' is not used [-Werror,-Wunused-private-field]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE bbq-audio-processing: same-family Gold-v4 anchor; verified_reference=true and passes the verified-set hard gates	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
62. Raw row 264 - fire-spread-responder robot-simulation / scheduling-simulation	BAD EVIDENCE misleading indentation; statement is not part of the previous 'if' [-Werror,-Wmisleading-indentation]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE pipeline-valve-actuator: same-family Gold-v4 direct target with verified_reference=true	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
63. Raw row 274 - rescue-frontier-explorer robot-simulation / scheduling-simulation	BAD EVIDENCE misleading indentation; statement is not part of the previous 'for' [-Werror,-Wmisleading-indentation]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE pipeline-valve-actuator: same-family Gold-v4 direct target with verified_reference=true	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.

64. Raw row 670 - solar-gap-interpolation-table table-pivot / grid-matrix-image	BAD EVIDENCE misleading indentation; statement is not part of the previous 'for' [-Werror,-Wmisleading-indentation]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE etl: verified table transformation with current executable Exercism tests	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.
65. Raw row 671 - student-letter-grade-table table-pivot / grid-matrix-image	BAD EVIDENCE misleading indentation; statement is not part of the previous 'for' [-Werror,-Wmisleading-indentation]	530 TREATMENT Excluded: this task ID is absent from the final 530 corpus. The raw target was not repaired or trained.
STRICT COMPILE FAILURE	GOOD REFERENCE etl: verified table transformation with current executable Exercism tests	TRAINING RULE A target must pass the pinned strict release compile before it can be a positive example; warnings-as-errors failures are terminal until repaired and rerun.

Appendix A3. Same-ID lineage conflicts (7)

These rows collide with a verified Gold ancestor under the same task ID. The final corpus retains the prior curated version and excludes the raw target.

66. Raw row 31 - avl-api-rate-limits balanced-search-tree / ordered-trees-indexes	BAD EVIDENCE The unverified raw target conflicts with the verified Gold ancestor under the same task ID. Answer similarity: 0.0711; Gold verified_reference=True; raw verification field is empty.	530 TREATMENT Ancestor retained: the final corpus uses the prior curated 401-row anchor version. Its prompt and answer hashes differ from the raw pair; the raw target is absent.
SAME-ID LINEAGE CONFLICT	GOOD REFERENCE avl-api-rate-limits: same-ID Gold-v4 anchor; verified_reference=true and selected from answer-blind verified signal-v3	TRAINING RULE A revision may replace a verified ancestor only after independent verification. Until then, keep one active lineage leaf and exclude the conflicting target.
67. Raw row 32 - avl-appointment-slots balanced-search-tree / ordered-trees-indexes	BAD EVIDENCE The unverified raw target conflicts with the verified Gold ancestor under the same task ID. Answer similarity: 0.0664; Gold verified_reference=True; raw verification field is empty.	530 TREATMENT Ancestor retained: the final corpus uses the prior curated 401-row anchor version. Its prompt and answer hashes differ from the raw pair; the raw target is absent.
SAME-ID LINEAGE CONFLICT	GOOD REFERENCE avl-appointment-slots: same-ID Gold-v4 anchor; verified_reference=true and selected from answer-blind verified signal-v3	TRAINING RULE A revision may replace a verified ancestor only after independent verification. Until then, keep one active lineage leaf and exclude the conflicting target.
68. Raw row 33 - avl-coupon-thresholds balanced-search-tree / ordered-trees-indexes	BAD EVIDENCE The unverified raw target conflicts with the verified Gold ancestor under the same task ID. Answer similarity: 0.0746; Gold verified_reference=True; raw verification field is empty.	530 TREATMENT Ancestor retained: the final corpus uses the prior curated 401-row anchor version. Its prompt and answer hashes differ from the raw pair; the raw target is absent.
SAME-ID LINEAGE CONFLICT	GOOD REFERENCE avl-coupon-thresholds: same-ID Gold-v4 anchor; verified_reference=true and selected from answer-blind verified signal-v3	TRAINING RULE A revision may replace a verified ancestor only after independent verification. Until then, keep one active lineage leaf and exclude the conflicting target.
69. Raw row 37 - avl-library-holds balanced-search-tree / ordered-trees-indexes	BAD EVIDENCE The unverified raw target conflicts with the verified Gold ancestor under the same task ID. Answer similarity: 0.0847; Gold verified_reference=True; raw verification field is empty.	530 TREATMENT Ancestor retained: the final corpus uses the prior curated 401-row anchor version. Its prompt and answer hashes differ from the raw pair; the raw target is absent.
SAME-ID LINEAGE CONFLICT	GOOD REFERENCE avl-library-holds: same-ID Gold-v4 anchor; verified_reference=true and selected from answer-blind verified signal-v3	TRAINING RULE A revision may replace a verified ancestor only after independent verification. Until then, keep one active lineage leaf and exclude the conflicting target.
70. Raw row 38 - avl-live-leaderboard balanced-search-tree / ordered-trees-indexes	BAD EVIDENCE The unverified raw target conflicts with the verified Gold ancestor under the same task ID. Answer similarity: 0.065; Gold verified_reference=True; raw verification field is empty.	530 TREATMENT Ancestor retained: the final corpus uses the prior curated 401-row anchor version. Its prompt and answer hashes differ from the raw pair; the raw target is absent.
SAME-ID LINEAGE CONFLICT	GOOD REFERENCE avl-live-leaderboard: same-ID Gold-v4 anchor; verified_reference=true and selected from answer-blind verified signal-v3	TRAINING RULE A revision may replace a verified ancestor only after independent verification. Until then, keep one active lineage leaf and exclude the conflicting target.

71. Raw row 125 - dll-book-shelf doubly-linked-list / queues-buffers-caches-lists	BAD EVIDENCE The unverified raw target conflicts with the verified Gold ancestor under the same task ID. Answer similarity: 0.0932; Gold verified_reference=True; raw verification field is empty.	530 TREATMENT Ancestor retained: the final corpus uses the prior curated 401-row anchor version. Its prompt and answer hashes differ from the raw pair; the raw target is absent.
SAME-ID LINEAGE CONFLICT	GOOD REFERENCE dll-book-shelf: same-ID Gold-v4 anchor; verified_reference=true and selected from answer-blind verified signal-v3	TRAINING RULE A revision may replace a verified ancestor only after independent verification. Until then, keep one active lineage leaf and exclude the conflicting target.
72. Raw row 404 - xor-card-game-turns xor-linked-list / queues-buffers-caches-lists	BAD EVIDENCE The unverified raw target conflicts with the verified Gold ancestor under the same task ID. Answer similarity: 0.1177; Gold verified_reference=True; raw verification field is empty.	530 TREATMENT Ancestor retained: the final corpus uses the prior curated 401-row anchor version. Its prompt and answer hashes differ from the raw pair; the raw target is absent.
SAME-ID LINEAGE CONFLICT	GOOD REFERENCE xor-card-game-turns: same-ID Gold-v4 anchor; verified_reference=true and selected from answer-blind verified signal-v3	TRAINING RULE A revision may replace a verified ancestor only after independent verification. Until then, keep one active lineage leaf and exclude the conflicting target.

Appendix B. All 129 verified complementary rows

Every row below shares a task ID with a review-only row in the raw 715 archive, but both its prompt and target changed before final selection. Every row is marked `verified_reference=true` and `executable_stage=passed`.

Task ID	Family	Semantic route	Mutation	Confidence
connector-grid-dilation	ascii-art-scaling	original-independent-judge	0.9375	0.98
flag-stripe-resampler	ascii-art-scaling	numeric-domain-rejudge	1.0000	0.98
flood-map-majority-downsample	ascii-art-scaling	numeric-domain-rejudge	1.0000	0.94
garden-tile-repeat	ascii-art-scaling	original-independent-judge	1.0000	0.98
label-runlength-expander	ascii-art-scaling	numeric-domain-rejudge	1.0000	0.99
minimap-viewport-zoom	ascii-art-scaling	original-independent-judge	1.0000	0.97
orchard-sparse-expander	ascii-art-scaling	numeric-domain-rejudge	1.0000	0.99
plate-coordinate-expander	ascii-art-scaling	numeric-domain-rejudge	0.8333	0.99
quilt-block-magnifier	ascii-art-scaling	numeric-domain-rejudge	0.9167	0.97
route-map-aspect-fit	ascii-art-scaling	original-independent-judge	1.0000	0.99
scale-archive-stamps	ascii-art-scaling	numeric-domain-rejudge	1.0000	0.96
seating-aisle-preserving-scale	ascii-art-scaling	original-independent-judge	1.0000	0.98
solar-palette-quantizer	ascii-art-scaling	numeric-domain-rejudge	1.0000	0.98
status-panel-decimator	ascii-art-scaling	original-independent-judge	1.0000	0.99
trail-polyline-raster-scale	ascii-art-scaling	targeted-repair-independent-rejudge	1.0000	0.97
wayfinding-letterbox	ascii-art-scaling	original-independent-judge	0.9091	0.99
crop-row-quarantine	flood-fill	targeted-repair-independent-rejudge	0.8750	0.99
oil-tide-wrap	flood-fill	targeted-repair-independent-rejudge	1.0000	0.98
downhill-patrol-zone-propagator	grid-ownership-mapping	targeted-repair-independent-rejudge	1.0000	0.98
gallery-room-label-audit	grid-ownership-mapping	targeted-repair-independent-rejudge	0.8750	0.98
garden-boundary-ledger	grid-ownership-mapping	targeted-repair-independent-rejudge	1.0000	0.99
marina-berth-ray-attribution	grid-ownership-mapping	original-independent-judge	1.0000	0.96
office-desk-voronoi-assignment	grid-ownership-mapping	targeted-repair-independent-rejudge	0.9333	0.93

Task ID	Family	Semantic route	Mutation	Confidence
orchard-row-quota-allocator	grid-ownership-mapping	targeted-repair-independent-rejudge	0.8667	0.97
parcel-component-perimeters	grid-ownership-mapping	original-independent-judge	1.0000	0.99
parking-permit-occupancy-audit	grid-ownership-mapping	original-independent-judge	1.0000	0.97
quay-concession-component-audit	grid-ownership-mapping	original-independent-judge	1.0000	0.99
rack-contiguous-allocation-audit	grid-ownership-mapping	original-independent-judge	1.0000	0.98
rail-platform-service-zones	grid-ownership-mapping	original-independent-judge	1.0000	0.98
river-junction-rights-propagator	grid-ownership-mapping	original-independent-judge	1.0000	0.99
shoreline-barrier-coverage	grid-ownership-mapping	original-independent-judge	1.0000	0.99
solar-string-serpentine-audit	grid-ownership-mapping	original-independent-judge	1.0000	0.98
warehouse-zone-capacity-counter	grid-ownership-mapping	targeted-repair-independent-rejudge	1.0000	0.97
gallery-door-region-graph	maze-to-graph	targeted-repair-independent-rejudge	1.0000	0.98
ski-elevation-dag	maze-to-graph	targeted-repair-independent-rejudge	1.0000	0.98
ridge-height-silhouette	pattern-printing	targeted-repair-independent-rejudge	1.0000	0.99
ski-chevron-route-sign	pattern-printing	targeted-repair-independent-rejudge	1.0000	0.99
striped-cone-stencil	pattern-printing	targeted-repair-independent-rejudge	1.0000	0.99
lane-blockage-run-auditor	run-length-image-encoding	targeted-repair-independent-rejudge	0.9333	0.98
radar-column-threshold-decoder	run-length-image-encoding	targeted-repair-independent-rejudge	0.9375	0.97
shelf-empty-span-index	run-length-image-encoding	targeted-repair-wave2-independent-xhigh-rejudge	1.0000	0.97
snow-map-delta-codec	run-length-image-encoding	original-independent-judge	0.8571	0.96
transparent-border-crop-decoder	run-length-image-encoding	targeted-repair-independent-rejudge	0.9412	0.95
vacancy-rectangle-run-index	run-length-image-encoding	targeted-repair-independent-rejudge	1.0000	0.95
defect-column-compressor	sparse-matrix-encoding	original-independent-judge	1.0000	0.99
failure-bipartite-index	sparse-matrix-encoding	original-independent-judge	1.0000	0.97
flood-shoreline-perimeter	sparse-matrix-encoding	original-independent-judge	1.0000	0.98
hotspot-row-sweep	sparse-matrix-encoding	numeric-domain-rejudge	1.0000	0.99
irrigation-gap-ledger	sparse-matrix-encoding	original-independent-judge	1.0000	0.99
warehouse-delta-coalescer	sparse-matrix-encoding	targeted-repair-independent-rejudge	1.0000	0.98

Task ID	Family	Semantic route	Mutation	Confidence
yield-row-dot-product	sparse-matrix-encoding	targeted-repair-independent-rejudge	0.9091	0.99
airport-delay-percentiles	table-pivot	numeric-domain-rejudge	1.0000	0.96
assay-weighted-mean-table	table-pivot	numeric-domain-rejudge	1.0000	0.96
branch-category-distinct-table	table-pivot	numeric-domain-rejudge	1.0000	0.98
clinic-state-transition-grid	table-pivot	original-independent-judge	1.0000	0.99
cohort-visit-retention	table-pivot	numeric-domain-rejudge	1.0000	0.96
defect-pareto-matrix	table-pivot	numeric-domain-rejudge	1.0000	0.99
depot-shortfall-table	table-pivot	targeted-repair-independent-rejudge	0.9000	0.98
exhibit-running-attendance	table-pivot	targeted-repair-independent-rejudge	1.0000	0.98
orchard-score-band-table	table-pivot	targeted-repair-independent-rejudge	1.0000	0.98
pivot-call-center	table-pivot	numeric-domain-rejudge	1.0000	0.99
river-unit-normalized-table	table-pivot	original-independent-judge	0.9333	0.98
room-occupancy-interval-grid	table-pivot	targeted-repair-independent-rejudge	1.0000	0.96
route-stop-cross-tab	table-pivot	original-independent-judge	1.0000	0.99
seasonal-yield-delta	table-pivot	numeric-domain-rejudge	1.0000	0.98
tiered-billing-pivot	table-pivot	numeric-domain-rejudge	1.0000	0.98
vendor-product-leader-table	table-pivot	numeric-domain-rejudge	1.0000	0.98
venue-capacity-grid	table-pivot	targeted-repair-independent-rejudge	1.0000	0.97
time-backup-cutover	clock-arithmetic	targeted-repair-independent-rejudge	0.9412	0.98
time-briefing-offset-board	clock-arithmetic	original-independent-judge	1.0000	0.99
time-irrigation-cycle	clock-arithmetic	targeted-repair-independent-rejudge	1.0000	0.99
time-medication-window	clock-arithmetic	original-independent-judge	1.0000	0.99
time-oven-program	clock-arithmetic	original-independent-judge	1.0000	0.98
time-overnight-roster	clock-arithmetic	targeted-repair-independent-rejudge	0.9375	0.98
time-parking-grace-audit	clock-arithmetic	original-independent-judge	1.0000	0.99
time-rail-transfer-checker	clock-arithmetic	targeted-repair-independent-rejudge	1.0000	0.99
time-school-bell-repair	clock-arithmetic	numeric-domain-rejudge	0.9333	0.96

Task ID	Family	Semantic route	Mutation	Confidence
timer-build-lease	countdown-timers	original-independent-judge	1.0000	0.92
timer-chess-round	countdown-timers	numeric-domain-rejudge	0.9333	0.97
timer-game-cooldown-registry	countdown-timers	original-independent-judge	1.0000	0.90
timer-incubator-checkpoints	countdown-timers	numeric-domain-rejudge	1.0000	0.97
timer-oven-safety-lock	countdown-timers	original-independent-judge	1.0000	0.96
timer-parking-credit	countdown-timers	numeric-domain-rejudge	0.9000	0.93
timer-study-session-budget	countdown-timers	original-independent-judge	1.0000	0.91
curfew-dispatch-selector	cross-midnight-intervals	original-independent-judge	0.9375	0.94
dock-booking-calendar	cross-midnight-intervals	original-independent-judge	1.0000	0.94
handoff-continuity-audit	cross-midnight-intervals	targeted-repair-independent-rejudge	0.9444	0.97
oven-capacity-scheduler	cross-midnight-intervals	original-independent-judge	1.0000	0.95
quiet-window-violations	cross-midnight-intervals	original-independent-judge	1.0000	0.97
dated-library-loan	date-difference	targeted-repair-independent-rejudge	1.0000	0.97
dated-project-burnup	date-difference	targeted-repair-independent-rejudge	1.0000	0.96
dated-retention-review	date-difference	targeted-repair-independent-rejudge	0.8824	0.94
dated-subscription-proration	date-difference	targeted-repair-independent-rejudge	0.9412	0.96
maintenance-event-ledger	duration-formatting	original-independent-judge	1.0000	0.99
parking-block-receipt	duration-formatting	original-independent-judge	1.0000	0.99
pipeline-stage-digest	duration-formatting	original-independent-judge	1.0000	0.99
ranked-lap-recap	duration-formatting	original-independent-judge	1.0000	0.99
reserve-discharge-forecast	duration-formatting	original-independent-judge	1.0000	0.99
split-budget-sla	duration-formatting	original-independent-judge	1.0000	0.99
turnback-air-plan	duration-formatting	original-independent-judge	1.0000	0.99
incident-budget-frontier	elapsed-time-accumulation	original-independent-judge	1.0000	0.98
machine-state-coverage-audit	elapsed-time-accumulation	targeted-repair-independent-rejudge	1.0000	0.93
podcast-attempt-selector	elapsed-time-accumulation	numeric-domain-rejudge	1.0000	0.97
service-outage-interval-union	elapsed-time-accumulation	targeted-repair-independent-rejudge	1.0000	0.97

Task ID	Family	Semantic route	Mutation	Confidence
construction-phase-network	future-date-calculations	targeted-repair-independent-rejudge	0.9444	0.95
orchard-treatment-window	future-date-calculations	original-independent-judge	1.0000	0.96
vaccine-eligibility-window	future-date-calculations	original-independent-judge	0.9375	0.90
warranty-service-calendar	future-date-calculations	targeted-repair-wave2-independent-xhigh-rejudge	0.7222	0.97
leap-archive-gap-index	leap-year-rule	original-independent-judge	1.0000	0.98
leap-benefit-apportionment	leap-year-rule	original-independent-judge	1.0000	0.94
leap-coverage-segments	leap-year-rule	original-independent-judge	0.9333	0.97
leap-maintenance-ledger	leap-year-rule	targeted-repair-independent-rejudge	0.9412	0.99
leap-policy-replay	leap-year-rule	original-independent-judge	1.0000	0.93
escalation-state-ledger	nth-and-final-occurrences	targeted-repair-independent-rejudge	0.9444	0.94
sla-window-crossings	nth-and-final-occurrences	targeted-repair-wave2-independent-xhigh-rejudge	1.0000	0.95
offset-auction-liquidity-intersection	offset-aware-range-overlap	original-independent-judge	0.9286	0.98
offset-border-slot-enumerator	offset-aware-range-overlap	original-independent-judge	1.0000	0.96
offset-deployment-quorum-sweep	offset-aware-range-overlap	numeric-domain-rejudge	1.0000	0.97
offset-observation-coverage-subtraction	offset-aware-range-overlap	original-independent-judge	0.9167	0.96
offset-rest-gap-compliance	offset-aware-range-overlap	targeted-repair-independent-rejudge	0.9375	0.95
offset-station-contact-weighted-schedule	offset-aware-range-overlap	targeted-repair-independent-rejudge	1.0000	0.95
offset-support-coverage-chain	offset-aware-range-overlap	original-independent-judge	0.8462	0.97
epoch-span-partitioner	overflow-safe-date-math	original-independent-judge	0.9444	0.96
preservation-policy-join	overflow-safe-date-math	original-independent-judge	1.0000	0.90
voucher-lifecycle-ledger	overflow-safe-date-math	targeted-repair-wave2-independent-xhigh-rejudge	0.9444	0.95
pack-address-lines	text-justification	original-independent-judge	1.0000	0.99
prevent-notice-widows	text-justification	original-independent-judge	1.0000	0.98
reflow-nested-quotes	text-justification	targeted-repair-wave2-independent-xhigh-rejudge	0.9333	0.96
segment-radio-cues	text-justification	targeted-repair-independent-rejudge	1.0000	0.95

Appendix C. Evidence limits and source ledger

LIMIT	The 715 audit does not prove that any individual row caused an evaluation outcome. It establishes row-level defects, missing verification, and mixture risks.
LIMIT	The 643 review rows were not labeled wrong. They lacked a supplied deterministic semantic oracle or current row-level verification receipt.
LIMIT	The final 530 is not a cleaned subset of the 715. It combines 401 inherited anchors with 129 changed and verified complements.
LIMIT	The final audit rechecked all 530 rows for structure and token length, and the 129 complements for strict compile/executable-stage status. The 401 anchors inherit their earlier evidence contract.
LIMIT	The completed SFT consumed 520 of 530 packaged rows because the existing loader executed only complete batches of 20.
LIMIT	The evaluation table is a receipt. It contains no per-row causal attribution or comparative judgment.

Frozen inputs and receipts

Artifact	Location	SHA-256 / receipt
Raw 715 archive	/Users/tokenbender/Downloads/aider-tasks-reverify-sft.zip	f90ff7c6c595b91f06b42c89e8d224f7f0d3786e5376ff17384aab59a13e9d38
Raw member	aider-tasks-reverify-sft/sft/train.jsonl	741d075a85f1f2f9c76b1ee03ea21e6355f9bb94e352563411d744be91c2bf9c
715 row decisions	/Users/tokenbender/Documents/Codex/2026-07-05/pie-slime-posttraining/artifacts/aider-tasks-regression-audit-20260720/row_decisions.jsonl	See audit SHA ledger
Final 530 train	/Users/tokenbender/Documents/Codex/2026-07-05/pie-slime-posttraining/artifacts/aider-cpp-complement-530/sft/train.jsonl	805aa59bbc936ee20687a293ef47d2fb9bcaee12419c6539ecd6180dfab02089
Original 401 archive	/Users/tokenbender/Downloads/train-401-dsa.zip	2efe714c454de7ba1c5fd523f5849b3c6c9af65e8e5ca5bf4cf438017fb1e03a
Combined-v5 archive	/Users/tokenbender/Downloads/aider-cpp-combined-v5-500.zip	3d2e3dc49a14e9b98c750520ccb34f76678e61b535e1b9c180b841e045842fbb
Final 530 run archive	/Users/tokenbender/Downloads/glm47-aider-complement-530-run-20260721.zip	c48edb54d3db7b659fc00e7a01cb32ce1352ac7d46a9ab0c0d788c8cb09006
Raw-concat eval receipt	/Users/tokenbender/Documents/Codex/2026-07-05/pie-slime-posttraining/artifacts/aider-cpp-raw-concat-v1-run-20260720/eval/glm47-aider-raw-concat-v1-sft-eval-20260720/run_receipt.json	Local receipt